

DataShield: Configurable Data Confidentiality and Integrity

Scott A. Carr, Mathias Payer

Purdue University

AsiaCCS 2017

Motivation



Heartbleed

- Missing bounds check in OpenSSL
- Leak private key
- Not prevented by deployed defenses
 - DEP, Stack Canaries, CFI

We need new tools to protect sensitive data

Introduction

Data Confidentiality and Integrity

- Some data are more sensitive
 - Worth paying overhead
 - Ex: stack canaries, DEP, CFI[1], CPI[2]
- Let the programmer choose
- Partial protection => Lower overhead
- Compiler inserts dynamic checks to protect sensitive data

1. Control-Flow Integrity. Abadi et al. CCS 2005
2. Code Pointer Integrity. Kuznetsov et al. OSDI 2014.

Our Assumptions

- Only low overhead is acceptable
 - 5-10% may be undetectable by the user [1]
- Have program source code
- Original program is buggy but benign
- Attackers exploit bugs to read/write unintended data

1. Everything You Want to Know About Pointer-Based Checking. Nagarakatte et al. SNAPL 2015.

What is Memory Safety?



Reads/writes through pointers R/W the object to which the pointer was most recently assigned

- Confidentiality \Leftrightarrow reads
- Integrity \Leftrightarrow writes

Spatial Memory Safety



```
i = 0;  
while (i <= 4) {  
    buf[i++] = 42;  
}
```



Why not Complete Memory Safety?

- Protecting all data is too costly
 - ~100% overhead [1,2]
- Overhead is a function of # of dynamic checks

1. SoftBound: Highly Compatible and Complete Memory Safety for C. Nagarakatte et al. PLDI 2009
2. CETS: Compiler-Enforced Temporal Safety for C. Nagarakatte et al. ISMM 2010.

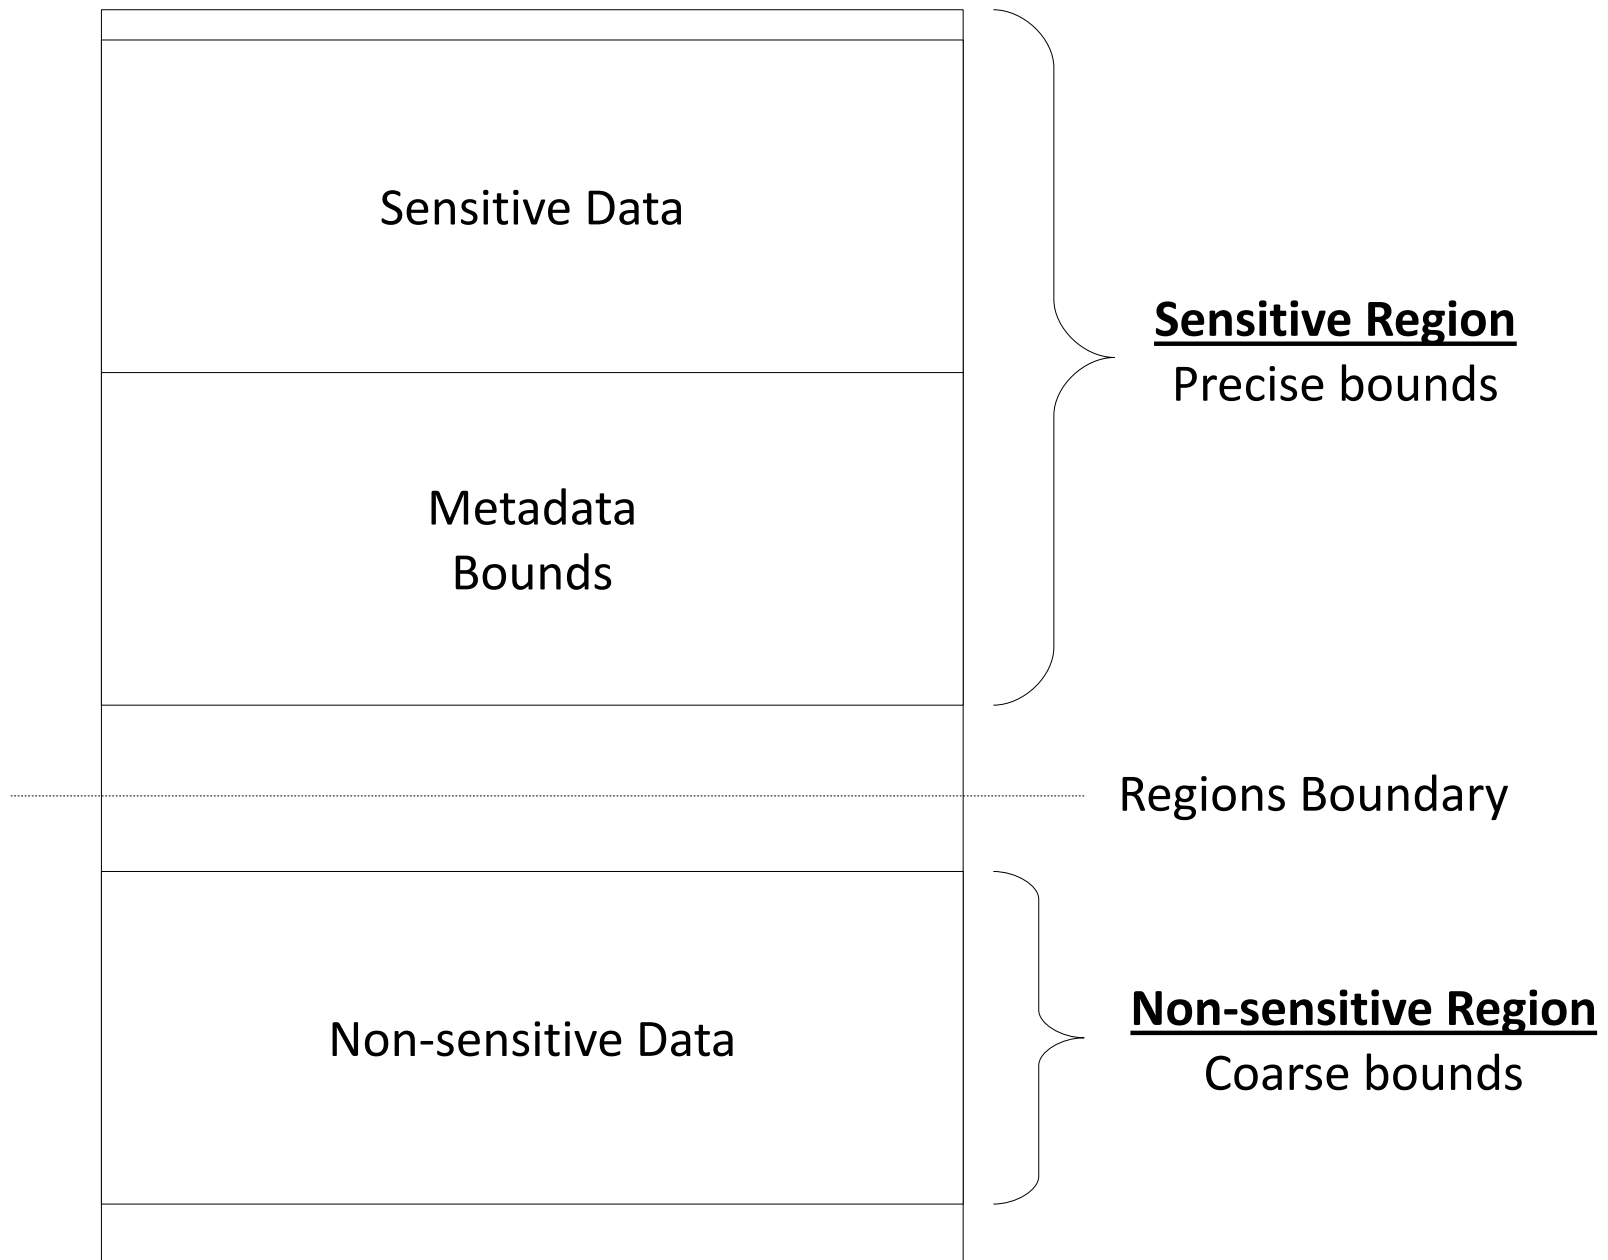
DCI Contributions

- New policy for protecting selected data
- Lower overhead than full memory safety
- Coarse bounds check implementations
- Security evaluation
 - Detects attacks in mbedTLS

Design

DCI Policy

- *Sensitive pointers* can only access the intended sensitive object
- *Non-sensitive pointers* can access any non-sensitive data
- Explicit data-flow between sensitive and non-sensitive objects is forbidden



Annotations

- Type based
- Marks type and members
- Cannot be cast away
- Mixed sensitivity **structs** not allowed

`__attribute__((annotate("sensitive")))`

Type Sensitivity

- All nested types have same sensitivity
- Pointers to sensitive types are sensitive
- Nesting a primitive type *P* does not make every *P* sensitive

```
struct S {  
    int x;  
    struct T *t;  
};  
  
struct T {  
    float z;  
    struct U *u;  
}  
  
struct U {  
    ...  
};
```

Implementation

Annotation Example (1)

```
struct foo {  
    char* name;  
    int x,y,z;  
};
```

```
__attribute__((annotate("sensitive"))) struct foo ignore;
```

Sensitive Allocation Example

```
struct foo* ptr = malloc(sizeof(struct foo));
```

Sensitive Allocation Example

```
struct foo* ptr = malloc(sizeof(struct foo));  
struct foo* ptr = sensitive_malloc(sizeof(struct foo));
```

```
table[&ptr].base = ptr;  
table[&ptr].end = ptr + sizeof(struct foo);
```

Sensitive Access Example

```
bounds = table[&ptr]
assert(bounds.base <= &(ptr->x));
assert(&(ptr->x) + sizeof(ptr.x) < bounds.end);
```

```
ptr->x = 5;
```

Non-sensitive Example

```
int* arr = malloc(sizeof(int)*100);
```

```
int* idx = arr+8;
```

```
*idx = 42;
```

Non-sensitive Example

```
int* arr = malloc(sizeof(int)*100);  
int* arr = non_sensitive_malloc(sizeof(int)*100);  
  
int* idx = arr+8;  
  
*idx = 42;  
  
int* idx_masked = idx & mask;  
*idx_masked = 42;
```

Implementation Overview

1. Compile time analysis

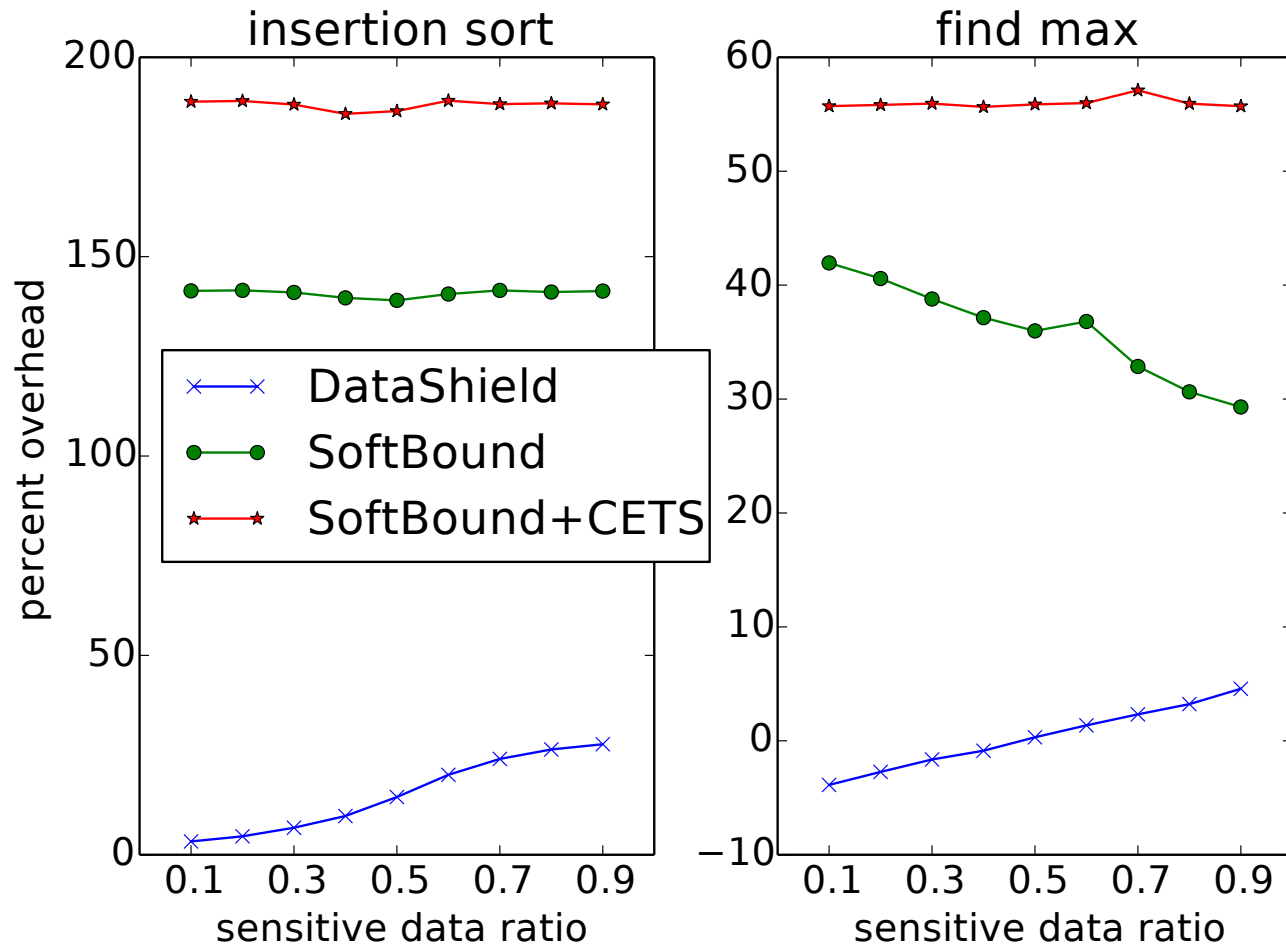
- Module-level analysis
- Inter-procedural type & context sensitive analysis

2. Runtime

- Separate sensitive/non-sensitive variables
- Heap, Stack, Global

Evaluation


Evaluation - Sensitivity Ratio



Evaluation – SPEC CPU2006

- No annotations added
- Sensitive regions bounds still enforced
- Measured overhead of code that does not access sensitive data

Evaluation – Coarse Bounds Checks



Security Evaluation

- CVE-2015-5291 from mbedTLS
- Malicious session ticket causes *buffer overflow*
- Proof of concept exploit publically available
- Compiled pre-patch version of mbedTLS
- Ran exploit
- Detected by DCI


Extra Slides


DCI Summary


- *Strong* protection for *sensitive data*
- *Weaker* (but lower overhead) protection for *non-sensitive*
- Lower overhead vs. complete memory safety
- Detects vulnerabilities in production software

<https://github.com/HexHive/datashield>

Extra Slides



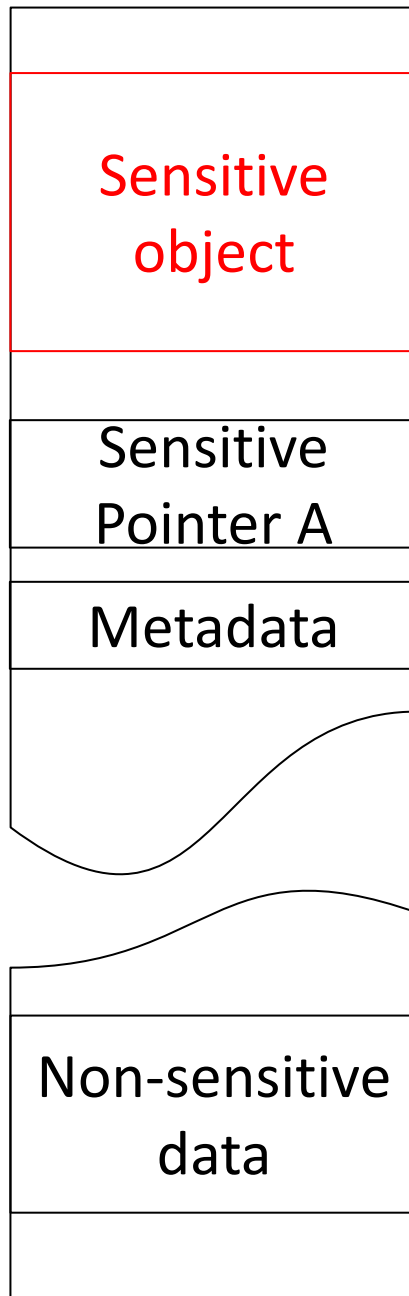




index: [base, end)
0x300: [0x200, 0x220)
0x210: [0x310, 0x320)
.
.
.
0xff: [0x0, 0xff)
0x0: [0x0, 0xff)



0x300



index: [base, end)
0x300: [0x200, 0x220)
0x210: [0x310, 0x320)
⋮
0xff: [0x0, 0xff)
0x0: [0x0, 0xff)



0x200
0x210
0x220

Standard Library Support

- Option 1: per-application lib
 - Rewrite specialized versions of each library function on-demand
 - Same analysis/rewriting as application
 - **Con:** Requires unique lib per application
 - **Pro:** Internal checks
- Option 2: drop in replacement lib
 - Make all library allocated data non-sensitive
 - Use wrappers and copies for sensitive
 - **Con:** No internal checks
 - **Pro:** Allows single, compatible lib

Evaluation - astar

- C++: 4,285 LoC
- Relaxed policy: separation mode
 - Primitive arithmetic does not propagate sensitivity
 - Reduces overhead 96% -> 9.12%
 - Reduce sensitive bounds checks by 10^6 times

Evaluation - mbedTLS

- C: 30,000 LoC
 - Instrument sample server and client
 - Annotate `ssl_context`
 - 35.7% overhead
-
- Challenge: sensitive data passed to callee through function pointer

Limitations

- Variadic arguments as sensitive
- Temporal metadata not implemented
- Region-based temporal protection if one sensitive type
 - Similar to Cling [1]

1. Cling: A Memory Allocation to Mitigate Dangling Pointers. P. Akritidis. USENIX Security 2010